



A study of reinforcement learning algorithms in simulated robotics scenarios

Alejandro Pajares Chirre

Masters Thesis submitted to the Faculty of AI at HS Fulda

Matriculation No: 1331534

Supervisor: Prof. Dr. Alexander Gepperth

Co-Supervisor: Prof. Dr. David James

Submitted on dd.mm.yyyy

Abstract

The world of robotics has seen significant advancements in recent years, and this is largely due to the integration of machine learning techniques. Robots are now able to learn from their surroundings, make decisions, and carry out tasks with minimal human intervention. Machine learning has enabled robots to interact with humans and perform tasks that were previously considered impossible. In particular, reinforcement learning (RL) is a type of machine learning that models how humans learn from sensory input and motor responses in response to rewards. RL is based on the idea that an agent interacts with an environment by taking actions and receiving feedback in the form of rewards or punishments. Q-learning is a popular algorithm used in RL to learn the optimal policy, i.e., the best sequence of actions to maximize reward, for an agent. This thesis focuses on the application of reinforcement learning (RL) and Q-learning algorithms in controlling the motion of three joints of a six-degree-of-freedom (6-DOF) robotic arm in a simulated environment. To apply the Q-learning algorithms, the problem needs to be modeled as a Markov Decision Process, and during the learning process, the exploration and exploitation rate need to be balanced. The robotic arm is modeled in Gazebo, and the control commands are sent using the Robot Operating System (ROS 2). The objective of the robotic arm is to touch the target. The RL algorithm learns to maximize the reward function, which is based on the current and previous distance between the target and one end of the robotic arm and the angles of the three joints being controlled. The experimental results demonstrate that RL and Q-learning algorithms can effectively control the motion of a robotic arm in a simulated environment. The robotic arm successfully learns to approach and touch the target.

Contents

| | |
|---|------------|
| List of Figures | III |
| List of Tables | III |
| 1 Introduction | 1 |
| 1.1 Context | 1 |
| 1.2 Problem statement | 1 |
| 1.3 Goals | 2 |
| 1.4 Related Work | 2 |
| 1.5 Contribution | 3 |
| 2 Foundations | 3 |
| 2.1 Topic 1 | 4 |
| 2.2 Topic 2 | 4 |
| 2.3 Topic 3 | 4 |
| 3 Implementation | 4 |
| 4 Experiments | 9 |
| 5 Discussion | 9 |
| 6 Conclusion | 9 |
| 7 Using LaTeX, erase this chapter later | 9 |
| 7.1 Mathematische Gleichungen | 9 |
| 7.2 Das ist eine Auflistung | 10 |
| 7.3 Das ist eine Bullet-Liste | 10 |
| 7.4 Eine Grafik bindet man so ein | 10 |
| 7.5 So schreibt man einen Algorithmus | 11 |
| 7.6 So gestaltet man eine Tabelle | 11 |
| 7.7 Interne Referenzen | 11 |
| 7.8 Textformatierung | 11 |
| 7.9 Zitieren | 11 |
| 7.10 Webquellen zitieren | 12 |
| 7.11 Literaturverzeichnis erstellen | 12 |
| A Code Snippets | 13 |
| B Thesis defence | 14 |

List of Figures

| | | |
|---|---|----|
| 1 | short caption | 5 |
| 2 | short caption | 6 |
| 3 | short caption | 6 |
| 4 | short caption | 7 |
| 5 | short caption | 8 |
| 6 | short caption | 8 |
| 7 | Logo der HAW Fulda | 10 |
| 8 | Normalerweise bindet man Snippets als Bilder ein... | 14 |

List of Tables

| | | |
|---|----------------------------|----|
| 1 | Beispielstabelle | 11 |
|---|----------------------------|----|

1 Introduction

Robotic systems are increasingly being used in a wide range of applications, from manufacturing and logistics to healthcare and entertainment. However, the effectiveness of these systems depends on their ability to perform tasks accurately and efficiently, which requires the development of intelligent control strategies that can adapt to different environments and objectives. Reinforcement learning offers a promising approach to achieve this goal, by enabling robots to learn from feedback and optimize their behavior based on the rewards received from the environment. Numerous companies have embraced the use of robotics to aid humans in tasks that are monotonous, physically demanding, or hazardous. Yet, acquiring a robot and hiring a robotic engineer to create a tailored solution for a particular task requires a significant investment of resources. The duties of a robot engineer include setting up communication, designing control scripts, computing coordinate transformations, and creating error-handling programs. Typically, a technician takes on the task of operating the robot on a daily basis, or the robot operates on its own. However, if the task requirements or processes change, it is difficult to modify the existing robotic solution to suit a new configuration or application without the assistance of a robot engineer, despite the significant resources invested in acquiring and developing it. Instead of relying on a robotic engineer to manually program a robot's operations for a new application, companies could employ deep reinforcement learning to train an intelligent agent to control the robot specifically for that application. This approach would enable the resources invested in robotics to be more adaptable and versatile, suitable for a broader range of applications and purposes.

This thesis is about reinforcement learning which is a type of machine learning that involves an agent learning from its interactions with an environment in order to maximize a reward signal over time [1]. It is a method of learning that involves trial and error, with the agent receiving feedback in the form of rewards or punishments for its actions. According to Kaelbling, Littman, and Moore [2], reinforcement learning can be defined as "a problem faced by an agent that learns behavior through trial-and-error interactions with a dynamic environment". Many different concepts and methodologies can be used to break down reinforcement learning. This work focuses on Deep Q learning and to study it a simulated robotic arm is trained to touch a can.

1.1 Context

What is the context of the presented work? If you work in a company, present the company first. Then the broad scientific or technological background in which the work is embedded should be presented and explained.

1.2 Problem statement

What is the problem that is being addressed or solved? Why is it important or beneficial to find a solution to this problem?

1.3 Goals

Here, give a list of bullet points of quantifiable goals of the presented work. These achievement of these goals will be shown in the experiments section. The list should have 3-4 entries. No blabla here, hard goals!

1.4 Related Work

Deep Q-learning has been successfully applied in various robotic control tasks, including manipulators and robotic arms. In this section, we summarize the results and contributions of six relevant papers on deep Q-learning for robotic arm control.

In [3], the authors proposed a method for robotic grasping using deep Q-networks (DQNs). The results showed that the proposed method outperforms other traditional grasping methods in terms of grasping success rate and efficiency.

In [4], the authors presented a framework for real-time control of a robotic arm using deep reinforcement learning. They used a DQN-based algorithm to learn the control policy for the robotic arm, and the results demonstrated that the proposed method can achieve accurate and robust control.

List works that have a similar goals, plus a short explanation (3-4 sentences at most) as to how they differ from your work. Do NOT make comparisons here (better than my work or similar), that happens in the discussion section.

Admissible related work is (in descending order of acceptability):

- Peer-reviewed scientific publications, ideally with a DOI. Use Google scholar for searching (GS can export BibTeX entries that you can copy into the .bib file of this project).
- White papers and publicly available documents without review, cite with title, URL and date of access. In addition, you need to submit the PDFs in electronic form.
- Web pages, especially for software projects (e.g., TensorFlow, nginx, react, Django). Cite via URL and date of access. A github/gitlab/etc link is acceptable as well. Nothing needs to be submitted electronically, but only use such sources of there is no other way.

Literature is cited like this: as shown in clemen1989combining, blablaba. Or: [?] has a similar scope in the domain of perverted numerical integration, however without considering the aspect of cupidity. Or: In [?], a study of perverted diagonal matrix perversions is presented. You need not include page numbers.

See also Kap. 7.9, 7.10.

1.5 Contribution

Here, you present a bullet list of your personal contributions to the topic of the thesis. For example:

- Implementation of a bash script that did not work and was hard to read
- Comparison of different implementations for matrix perversion
- Implementation of a web service that provides jokes about professors via a ReST API.

2 Foundations

The targeted group are computer scientists with at least a Bachelor's degree. Here, you explain aspects that go beyond what this group would not usually know. For example:

- Specialized libraries and their use
- Complex concepts of the chosen programming language
- Basics of machine learning, or neural networks, or both
- Description of used databases or datasets
- rviz
- rqt
- urdf file
- sdf file
- stl file

In subsequent chapters, you can reference this one to avoid having to explain everything over and over again. This means that you just include things here that are necessary for the understanding of later chapters, nothing more.

2.1 Topic 1

2.2 Topic 2

2.3 Topic 3

3 Implementation

The simulated scenario is implemented using Gazebo and ROS 2 libraries, which are described in chapters AAA and BBB respectively. The robotic arm kuka kr210 simulation was taken from <https://github.com/udacity/RoboND-Kinematics-Project>, this was done by taking the urdf, dae, and stl files containing the specifications of the robotic arm and placing them in our ROS 2 package. The robotic arm consists of six joints and seven links that connect them. The six joints of the KR210 are numbered J1 to J6, and they allow the robot to move in various directions and orientations.

- J1: The first joint is the base joint, which allows the robot to rotate horizontally around its vertical axis.
- J2: The second joint is the shoulder joint, which allows the robot to lift and lower its arm vertically.
- J3: The third joint is the elbow joint, which allows the robot to bend its arm vertically.
- J4: The fourth joint is the wrist roll joint, which allows the robot to rotate its wrist around its vertical axis.
- J5: The fifth joint is the wrist pitch joint, which allows the robot to tilt its wrist up and down.
- J6: The sixth joint is the wrist yaw joint, which allows the robot to rotate its wrist horizontally.

The arm also has links that connect the joints, including the base, lower arm, upper arm, wrist, and end-effector. These links are designed to provide strength and rigidity to the robot arm while allowing for smooth and precise movement. These links are:

- Base Link: This is the fixed part of the robot that is attached to the ground.
- Link 1: This is the first link that connects the base to joint 1.
- Link 2: This link connects joint 1 to joint 2.
- Link 3: This link connects joint 2 to joint 3.
- Link 4: This link connects joint 3 to joint 4.

3 Implementation

- Link 5: This link connects joint 4 to joint 5.
- End Effector (Link 6): This is the final link that connects to the tool or object being manipulated.

A Bumper sensor was added to links 4, 5, and 6 to detect collisions with these links (As shown in the code snippet).

The object the robotic arm has to learn to touch is a can taken from gazebo models (<http://models.gazebo-sim.org/beer>). A world file is created containing the can.




Figure 1: long caption

3 Implementation




Figure 2: long caption




Figure 3: long caption




Figure 4: long caption




Figure 5: long caption




Figure 6: long caption

Should refer, where possible, to the preceding chapter, e.e.: Singular value decom-

7 Using LaTeX, erase this chapter later

position of the matrix Σ is conducted as explained in Sec. ?? using the *lapack* library (see Sec. 2.2).

For software development: what is the logic of the developed code, which of it was done by yourself? Sequence diagrams or UML are good tools here.

Please give code snippets only if they take up less than 0.25 pages, and only if it is unavoidable. Longer snippets go to the appendix and are referenced like this: see App. 8.

4 Experiments

Show here that the goals from the introduction were achieved (or not achieved), you need at least one experiment per goal. Use screenshots, diagrams, plots, photos, etc. as necessary.

5 Discussion

2-3 pages are a good idea here. Picks up goals from the introduction (see 1.3) and experiments (see 4) and explains what was achieved and what was not (and why not in this case). Compares results with results from related work, see Sec. ?. Draws a preliminary conclusion for the whole thesis.

6 Conclusion

Give an executive summary for important decision makers here, as well as an outlook (what would you do if you had another 3 months). 2-3 pages are ok here.

7 Using LaTeX, erase this chapter later

I was too lazy to translate this, it will be translated later. But I believe the ideas are clear!

7.1 Mathematische Gleichungen

Eine mehrzeilige Gleichung sieht so aus (die Symbole nach den und-Zeichen werden untereinander gesetzt). Die nonnumber-Befehle verhindern dass die Gleichung nummertiert wird (Geschmackssache, ist nie falsch wenn eine Gleichung nummeriert ist). Aber: eine

Gleichung auf die man referenziert (also die ein Label hat), muss nummeriert sein!

$$\begin{aligned} A &= \sum_{i=1}^N x_i \\ B &= \frac{\pi}{2} \end{aligned} \tag{1}$$

Eine inline-Gleichung: $x = 45b + \frac{2}{3}\pi$. Der Text geht weiter! Auf inline-Gleichungen kann man keine Referenzen erstellen.

7.2 Das ist eine Auflistung

1. Element 1
2. Element 2

7.3 Das ist eine Bullet-Liste

- Element 1
- Element 2

7.4 Eine Grafik bindet man so ein

Zulässige Formate sind generell eps, pdf und png.

Hochschule Fulda
University of Applied Sciences



Figure 7: Logo der HAW Fulda

7.5 So schreibt man einen Algorithmus

Algorithm 1: How to write algorithms

Data: this text

Result: how to write algorithm

initialization;

while *not at end of this document* **do**

 read current;

if *understand* **then**

 go to next section;

 current section becomes this one;

else

 go back to the beginning of current section;

end if

end while

7.6 So gestaltet man eine Tabelle

Table 1: Beispielstabelle

| A | B | C |
|---|----------|-------|
| D | per gram | 11.65 |
| | each | 1.01 |
| E | stuffed | 32.54 |
| F | stuffed | 73.23 |
| G | frozen | 8.39 |

7.7 Interne Referenzen

So wird ein Kapitel oder Unterkapitel referenziert: Kap. 1, Kap. 7.10. Auf Gleichungen bezieht man sich so: Wie in Gl. (1) gezeigt, sehen Gleichungen in der Regel gut aus. Auf Abb. 7 bezieht man sich so. Auf Tab. 1 referenziert man so. Algorithmen sind analog: siehe Alg. 1. Generell kann man alles zitieren was ein Label hat.

7.8 Textformatierung

So wird dick geschrieben und *so kursiv*.

7.9 Zitieren

Generell zitiert man so: wie in [?] gezeigt, blabla. Für jedes zitierte Werk ist ein BibTex-Eintrag nötig! Eine gute Quelle ist Google Scholar!!

7.10 Webquellen zitieren

So wird eine Webquelle zitiert: [5], siehe auch den Eintrag im BibTeX-File. Wichtig: für jede Web-Quelle ein BibTeX-Eintrag! Wenn Sie das auf die hier gezeigte Art machen, werden URLs (fast) automatisch getrennt. Kontrollieren Sie trotzdem die Literaturliste, es kann sein dass das nicht immer funktioniert.

7.11 Literaturverzeichnis erstellen

Hierzu müssen BibTeX-Einträge in die Datei literatur.bib eingefügt werden. Die BibTeX-Keys sind jeweils Argumente für die cite-Kommandos! Wenn Sie literatur.bib ändern müssen Sie alles mindestens 5x compilieren: 3x mit latex, 1x mit BibTeX und dann noch 2x mit LaTeX (in der Reihengfolge). Am besten Sie machen ein Skript dafür!

References

- [1] R. S. Sutton and A. G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2018.
- [2] M. L. Littman L. P. Kaelbling and A. R. Moore. Reinforcement learning: A survey. *Journal of Artificial Intelligence Research*, 4:237–285, 1996.
- [3] Sergey Levine, Peter Pastor, Alex Krizhevsky, Julian Ibarz, and Deirdre Quillen. Learning hand-eye coordination for robotic grasping with deep learning and large-scale data collection. *The International Journal of Robotics Research*, 37(4-5):421–436, 2018.
- [4] Shixiang Gu, Ethan Holly, Timothy Lillicrap, and Sergey Levine. Deep reinforcement learning for robotic manipulation with asynchronous off-policy updates. In *2017 IEEE International Conference on Robotics and Automation (ICRA)*, page 3389–3396. IEEE Press, 2017.
- [5] RStudio. Welcome to shiny. <https://shiny.rstudio.com/tutorial/written-tutorial/lesson1/>.

A Code Snippets

```
<gazebo reference="link_4">
  <!-- contact sensor -->
  <sensor name="end_effector_sensor" type="contact">
    <selfCollide>true</selfCollide>
    <alwaysOn>true</alwaysOn>
    <update_rate>500</update_rate>
    <contact>
      <collision>link_4_collision</collision>
    </contact>
  <!-- gazebo plugin -->
  <plugin name="gazebo_ros_bumper_sensor" filename="libgazebo_ros_bumper.so">
    <ros>
      <namespace>contact_sensor</namespace>
      <remapping>bumper_states:=bumper_link_4</remapping>
    </ros>
    <frame_name>link_4</frame_name>
  </plugin>
</sensor>
</gazebo>
```

Listing 1: Example XML listing using minted

Figure 8: Normalerweise bindet man Snippets als Bilder ein...

B Thesis defence

The defence is 15/20 minutes for Bachelor/Master, followed by questions and a discussion. Both examiners are present, and you can invite external persons since defences are generally public.

Targetted group are non-computer scientists, e.g., from higher management, NOT the examiners. Means that at least $\frac{1}{3}$ if the presentation is introduction/context/problem statement. You should re-use text/images/graphs/etc from the corresponding chapters here!

1 Slide per minute is a good guideline. If you can guess that some questions are going to be asked anyway, prepare some slides specifically for these questions, makes a good impression, and you can show them in the discussion time, not during the 15 minutes of the presentation.

Defences are not graded, you can only pass or not pass.

Students are responsible for finding dates for the defence and coordinating this with both supervisors.

Some common advice is:

- Speak slowly and loadly
- If you do not have enough time left for all slides, leave some out rather than rushing through all of them!!
- Slide numbers!
- In presence: be there 10 minutes ahead of time to check projectors etc. Makes a very bad impression if this is not working. Same for online presentations: be there 5 minutes ahead of time to verify screen sharing works.
- do not read text from the slides. These should contains key words only, and you explain the rest in free presentation
- Defences can by all means be online, more convenient for companies
- in presence: always carry a USB key with a PDF of your slides. If you have to use another PC than yours, PowerPoint slides may look very differently (fonts, page setup etc.)
- No-Go: spelling errors on slides!!!
- Do not use animations, they may not work in an online setting